



# Introduction to Terraform (2d)

Michael Bright / @mjbright Consulting



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright

@mjbright CONSULTING

# Introduction

# Introduction




# Introduction

Terraform is a tool for provisioning computing infrastructure

It's defining feature is managing “Infrastructure as Code” (IaC)

This means that we define the “desired state” of the infrastructure as code, in configuration files.




# Introduction

“Infrastructure as Code” (IaC) can be achieved through

- Adhoc scripting                      Shell, Python, Perl, ...
- Config management tools      Ansible, Chef, Puppet, Salt, ...
- Server templating tools      Docker, Packer, Vagrant
- Orchestration tools              Kubernetes, Mesos, Docker Swarm, ECS
- Provisioning tools               Terraform, AWS CloudFormation, OpenStack Heat

# Introduction

We tell terraform to “apply” the configuration files to create the appropriate resources.



If we tell terraform to “re-apply” the configuration then no further changes will be made – unless the configuration state has drifted.

This means that Terraform is idempotent.

# Introduction

“Infrastructure as Code” (IaC) brings

- Automation, repeatability
- Documentation
- Version Control, Audit trail
- Validation, Testability
- Reuse

# Introduction

## Terraform

- supports many cloud providers, i.e. it is provider agnostic.
- supports many resources for each provider.
- define resources as code in Terraform templates.



# Terraform Configurations

# Example Terraform configuration

Terraform configurations are written in the Hashicorp Configuration Language (HCL v2) & specify a provider & several resources

The Terraform CLI is a general engine for evaluating & applying configurations, using provider plugins which define & manage a set of resource types.

```
provider 'aws' {  
    region = 'us-west-1'  
}  
  
resource 'aws_instance' 'example' {  
    ami           = 'ami-408c7f28'  
    Instance_type = 't2.micro'  
    Tags = { Name = 'terraform_example' }  
}
```

# Example Terraform configuration

The Terraform language is used to declare resources specific to a provider.

```
resource "aws_vpc" "main" {  
    cidr_block = var.base_cidr_block  
}
```

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body: contains arguments or nested blocks  
    <IDENTIFIER> = <EXPRESSION> # Argument  
}
```

<https://www.terraform.io/docs/configuration/index.html>

# Example Terraform configuration

Groups of resources can be gathered into modules

A *Terraform configuration* consists of a *root module*, where evaluation begins, along with a tree of child modules created when one module calls another.

<https://www.terraform.io/docs/configuration/index.html>

# Terraform Providers

It is required to have a provider defined for your terraform configuration

We will be using AWS, but Terraform supports many providers

# Terraform Providers

## AWS

Alibaba Cloud, Azure, DigitalOcean, Exoscale, Google Cloud Platform, Heroku, Oracle Cloud, OVH, Packet, 1&1, Spotinst

Terraform Cloud, Vault, Nomad, Consul

<https://www.terraform.io/docs/providers/index.html>

GitHub, GitLab, Bitbucket,

Kubernetes, Helm, Docker, OpenStack

PostgreSQL, MySQL, MongoDB Atlas

VMware NSX-T, vCloud , vSphere

Chef, Cobbler, Datadog, DNS, HTTP, Local, TLS

# Terraform – Provider “AWS” Services

ACM PCA API Gateway Application Autoscaling AppMesh AppSync Athena  
Autoscaling Backup Batch Budgets Cloud9 CloudFormation CloudFront  
CloudHSM v2 CloudTrail CloudWatch CodeBuild CodeCommit CodeDeploy  
CodePipeline Cognito Config Cost & Usage Report Data Lifecycle Manager  
(DLM) Database Migration Service (DMS) DataPipeline DataSync Device Farm  
Directory Service Direct Connect DynamoDB Accelerator (DAX) DocumentDB  
**EC2** ECR ECS EFS EKS ElastiCache Elastic Beanstalk **Elastic Load Balancing v2**  
**(ALB/NLB)** Elastic Map Reduce (EMR) ElasticSearch Elastic Transcoder  
Firewall Manager (FMS) File System (FSx) Gamelift Glacier Global Accelerator  
Glue GuardDuty **IAM** IoT Inspector Kinesis Kinesis Firehose KMS **Lambda**  
License Manager Lightsail Macie MQ MediaPackage MediaStore Managed  
Streaming for Kafka (MSK) Neptune OpsWorks Organizations Pinpoint Pricing  
QuickSight RAM **RDS** Redshift Resource Groups **Route53** Resolver **S3**  
Sagemaker Secrets Manager Security Hub SES Service Catalog Service  
Discovery Service Quotas Shield SimpleDB SNS SQS SSM Step Function (SFN)  
Storage Gateway SWF Transfer **VPC** WAF WAF Regional WorkLink  
WorkSpaces XRay

# Terraform Installation



# Terraform Installation

Terraform is a single static go binary available for macOS, Windows, Linux, downloadable from:

- <https://www.terraform.io/downloads.html>
- <https://github.com/hashicorp/terraform/releases>

On Unix-like OS, simply put it in your PATH, e.g.

```
mv terraform ~/bin/
```

may require:

```
export PATH=$PATH:~/bin
```

<https://github.com/hashicorp/terraform>

# Terraform Installation - Version

```
> terraform version  
Terraform v0.12.12  
+ provider.aws v2.33.0  
+ provider.template v2.1.2
```

Your version of Terraform is out of date!

The latest version is 0.12.13.

You can update by downloading from [www.terraform.io/downloads.html](http://www.terraform.io/downloads.html)

# Lab 0 – Installing Terraform

In this initial lab we will simply install the Terraform static binary for amd64 architecture and add this to our PATH

# Quiz

What does it mean to say that Terraform is idempotent ?

What are some advantages of “*Infrastructure as Code*” ?

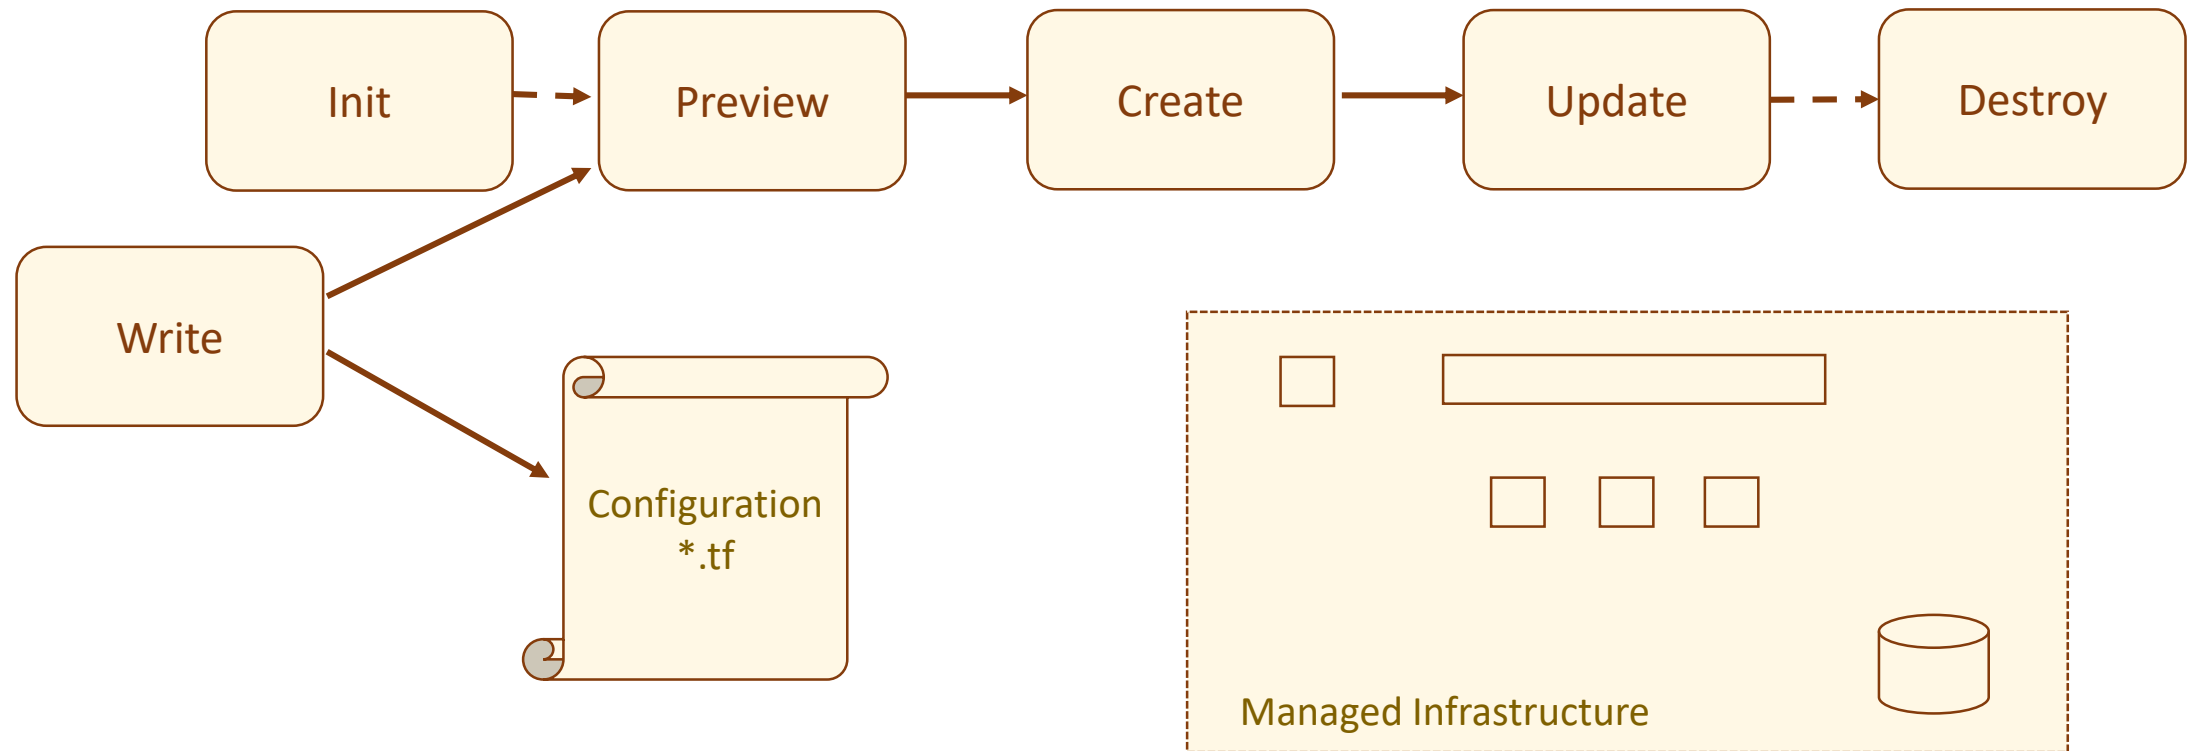
Is Terraform an Open Source project ?

What company is behind Terraform ?

# Terraform Workflow

# Terraform Workflow

Use of Terraform is according to the workflow




<https://www.terraform.io/docs/commands/index.html>

# Terraform Workflow

Use of Terraform is according to the workflow

- terraform init - download and initialize specified providers
- terraform plan - parse and check config
- terraform apply - create or update all resources
- terraform destroy - destroy all resources



<https://www.terraform.io/docs/commands/index.html>

# Terraform Workflow

<b>apply</b>	Builds or changes infrastructure	<b>plan</b>	Generate and show an execution plan
console	Interactive console for Terraform interpolations	providers	Prints a tree of the providers used in the config
<b>destroy</b>	Destroy Terraform-managed infrastructure	refresh	Update local state file against real resources
env	Workspace management	show	Inspect Terraform state or plan
fmt	Rewrites config files to canonical format	taint	Manually mark a resource for recreation
get	Download install modules for the config	untaint	Manually unmark a resource as tainted
graph	Create a visual graph of Terraform resources	validate	Validates the Terraform files
import	Import existing infrastructure into Terraform	version	Prints the Terraform version
<b>init</b>	Initialize a Terraform working directory	workspace	Workspace management
output	Read an output from a state file		



# Terraform Workflow - Get

The get command is used to download & update modules referenced in the root module.

```
terraform get [-update] [<dir>]
```

-update - reload modules even if already present.

dir - Sets the path of the root module.

Modules download to .terraform unless overwritten by TF\_DATA\_DIR.

e.g. TF\_DATA\_DIR=~/.dot.terraform # to avoid version control (can be large)

It is safe to run this command multiple times

<https://www.terraform.io/docs/commands/get.html>

# Terraform Workflow - Init

The init command is the first command which should be run to initialize a working directory of configuration files

```
terraform init [<options>] [<dir>]
```

Also downloads modules referenced in root configuration

It is safe to run this command multiple times

<https://www.terraform.io/docs/commands/init.html>

# Terraform Workflow - Init

```
> terraform init
```

```
Initializing the backend...
```

```
Initializing provider plugins...
```

- Checking for available provider plugins...
- Downloading plugin for provider "aws" (hashicorp/aws) 2.34.0...

```
The following providers do not have any version constraints in configuration, so the latest version was installed.
```

```
To prevent automatic upgrades to new major versions that may contain breaking changes, it is recommended to add version = "..." constraints to the corresponding provider blocks in configuration, with the constraint strings suggested below.
```

```
* provider.aws: version = "~> 2.34"
```

```
Terraform has been successfully initialized!
```

```
...
```

```
https://www.terraform.io/docs/commands/init.html
```

# Terraform Workflow - plan


The *plan* command creates an execution plan.

```
terraform plan [-out]
```

Refreshes it's view of the *actual* state, then determines what actions are necessary to achieve the *desired* state

- Checks the configuration file
- Show what will be done when the configuration is applied

Convenient to **preview** the effect of applying a configuration



<https://www.terraform.io/docs/commands/plan.html>

# Terraform Workflow - plan

```
> terraform plan
+ aws_instance.example
  ami:          "" => "ami-408c7f28"
  instance_type: "" => "t2.micro"
  key_name:     "" => "<computed>"
  private_ip:   "" => "<computed>"
  public_ip:    "" => "<computed>"
```

```
Plan: 1 to add, 0 to change, 0 to destroy
```

# Terraform Workflow - validate

You can also validate the syntax of files using the “*terraform validate*” command


```
> terraform validate  
Success! The configuration is valid.
```

# Terraform Workflow - apply

*The apply* command applies the changes required to reach the desired state of the configuration

```
terraform apply [<options>]
```

**Creates** the required resources, e.g. VM instances



<https://www.terraform.io/docs/commands/apply.html>

# Terraform Workflow - apply

```
> terraform apply
+ aws_instance.example: Creating...
  ami:                "" => "ami-408c7f28"
  instance_type:      "" => "t2.micro"
  key_name:           "" => "<computed>"
  private_ip:         "" => "<computed>"
  public_ip:          "" => "<computed>"
aws_instance.example: Creation complete
```

Apply complete! Resources: **1 added**, 0 changed, 0 destroyed.



# Terraform Workflow - apply


## Warning!

It is not guaranteed that if “*terraform plan*” returns no errors that the “*terraform apply*” will necessarily work!

If you attempt to do something that the provider doesn't support or allow, then the apply will fail even though the plan output doesn't report any errors!

# Terraform Workflow - apply

The configuration may be updated and “*terraform apply*” run again to **update** the resources to the new *desired* state




# Terraform Workflow - destroy

To destroy resources, use the ***destroy*** command.

```
> terraform destroy
aws_instance.example: Refreshing state... (ID: i-f3d58c70)
aws_elb.example: Refreshing state... (ID: example)
aws_elb.example: Destroying...
aws_elb.example: Destruction complete
aws_instance.example: Destroying...
aws_instance.example: Destruction complete

Apply complete! Resources: 0 added, 0 changed, 2 destroyed.
```



# Terraform – AWS Setup

You will need to export your AWS keys as environment variables.  
On Mac or Linux, do this with these commands:

```
export AWS_ACCESS_KEY_ID=<access_key>  
export AWS_SECRET_ACCESS_KEY=<secret_key>
```

On Windows, use set instead of export.

AWS\_DEFAULT\_REGION may be exported to provide a default region value

# Lab 1 – Introduction to Terraform

In this lab we will create an initial configuration to launch a single AWS EC2 virtual machine instance.

We will preview, then create this configuration following the basic Terraform workflow

We will then enable ssh access to show that we can have ssh access to the instance

We will use the “terraform graph” command to create a graphical representation of the Terraform resources we created

# Quiz

What does “*terraform init*” do ?

When was Terraform created ?

What happens when you type “*terraform apply*” ?

# Terraform Templates

# Parameterizing Terraform templates

We can parameterize our templates using *variables*.

Description, default and type attributes are optional.

```
variable "name" {  
    Description = "The name of the EC2 instance"  
}
```



# Parameterizing Terraform templates

Note: the use of the “\${}” syntax to interpolate var.name

```
variable "name" {
    description = "The name of the EC2 instance"
}

resource "aws_instance" "example" {
    ami = "ami-408c7f28"
    instance_type = "t2.micro"
    tags {
        name = var.name                # simplified terraform 0.12 syntax
        info = "${var.name} info"     # Required syntax for complex interpolation
    }
}
```

# Submitting a value

Note: in the previous example we didn't set a value for the variable, only its description.

When run with *“plan”* or *“apply”* we will be prompted for the value

```
> terraform plan
The name of the EC2 instance
Enter a value: foo

~ aws_instance.example
  Tags.Name: "terraform-example" => "foo"
```

# Submitting a value

We can also pass values to variables by

- using the --var parameter on the command line or
- Setting environment variable TF\_VAR\_foo

```
> terraform apply --var name=foo
aws_instance.example: Refreshing state...
aws_instance.example: Modifying...
tags.Name: "terraform-example" => "foo"
aws_instance.example: Modifications complete.

Apply complete! Resources: 0 added, 1 changed, 0 destroyed.
```

# Submitting a value

Variables can be set:

- By their “*default*” value field
- By the --var command-line option
- By the --var-file <file> command-line option
- By environment variables of the form TF\_VAR\_variable for “variable”
- From terraform.tfvars or \*.auto.tfvars files if present
- If a variable is unassigned at plan/apply time - Terraform will prompt for a value

# Submitting a value

## Precedence for variables

- If the default value for the variable is defined it is used
- If TF\_VAR\_<var> is set e.g. TF\_VAR\_varname this overrides
- If --var var=val is passed as an argument this takes precedence

# Submitting a value

Precedence for the region attribute/variables

- If region variable is unset & AWS\_DEFAULT\_REGION is set, it is used
- If the default value for region is defined it is used
- If TF\_VAR\_region is set, this overrides
- If --var var=val is passed as an argument this takes precedence

# Lab 2 – Terraform Variables

In this lab we will see how we can parametrize the configuration template using variables

# Quiz

When would you use “\${variable}” syntax, and when not ?

What are the 5 ways you can define variables for Terraform ?



# Terraform Types

# Terraform Types

Terraform supports different types of variables

- Numeric
- Booleans
- Lists
- Maps
- Sets
- any, null

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Types: Simple Types

Terraform supports different types of variables

Simple Types:

- string: a sequence of Unicode characters e.g. "hello"
- number: a numeric value. Integer or floating point
- boolean: true or false. bool values can be used in conditional logic

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Types: Simple Types

The standard way to create a variable is by simple declaration:

```
variable "mystr" {  
    default = "some_value"  
    description = "just a string"  
}
```

Both the default and descriptions are optional

Terraform recognizes the type as "string" in this case

```
variable "mynum" {  
    default = 23.7  
    description = "just a number"  
}  
  
variable "mybool" { default = true }
```

# Terraform Types: Strings

Simple strings are enclosed in “ ” quotes:

```
myvar = "a string"
```

Terraform supports multiline strings:

```
template = <<-EOF
    #!/bin/bash
    run-microservice.sh
    EOF
}
```

# Terraform Types: Number arithmetic

Terraform also supports basic arithmetic operations.

- + - \* / % for integers
- + - \* / for floats.

```
variable "myvar" {  
    default = 1 + 1  
    description = "one plus one"  
}  
  
output "even_1_1_3_mod2" { value = (1 + 1 + 3) % 2 == 0 } # false  
output "even_1_1_4_mod2" { value = (1 + 1 + 4) % 2 == 0 } # true
```

<https://www.terraform.io/docs/configuration/expressions.html#arithmetic-and-logical-operators>

# Terraform Types: Boolean

Terraform also supports booleans

```
variable mytruth {  
    default = true  
    description = "is this really true?"  
}  
  
output mytruth          { value = var.mytruth }  
output is_mytruth_true { value = (var.mytruth == true) }  
output bool_true       { value = true }  
output not_bool_true   { value = !true }
```

Outputs:

```
_true          = true  
is_mytruth_true = true  
mytruth        = true  
not_bool_true  = false
```

<https://www.terraform.io/docs/configuration/expressions.html#arithmetic-and-logical-operators>

# Terraform Types : Simple Types

## Conversion of Primitive Types

Terraform automatically converts number & bool values to string values as needed, & vice-versa as long as the string contains a valid representation of a number or boolean value

- true converts to "true", and vice-versa
- false converts to "false", and vice-versa
- 15 converts to "15", and vice-versa

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>



# Terraform Types : Collections

## Complex Types: Collections

- list: sequence of values – all the same type - indexed from zero.
- map: collection key/value pairs, where keys are strings
- set: collection of unique values

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Types: Collections - Lists

Terraform supports lists:

```
variable "mylist" {  
    type = "list" # not needed  
    default = ["foo", "bar", "baz"]  
}
```

Note that type is optional.

Terraform can figure out that it's a list from the default syntax.

# Terraform Types: Collections - Lists

We can access individual elements in the list by using the *element()* function, or with other syntaxes:

```
variable mylist {  
    type = list(number) # unneeded  
    default = [1,2,3]    # All elements must be of same type  
}
```

```
output elem0 { value = element(var.mylist, 0) }  
output elem1 { value = element(var.mylist, 1) }  
output elem1a { value = var.mylist.1 }  
output elem1b { value = var.mylist[1] }
```

```
output list { value = var.mylist }  
output lista { value = var.mylist.* }
```

Outputs:

```
elem0 = 1  
elem1 = 2  
elem1a = 2  
elem1b = 2
```

```
list = [ 1, 2, 3, ]  
lista = [ 1, 2, 3, ]
```

# Terraform Types: Complex - Object

## Complex Types: Structural Types

- object: collection of named attributes that each have their own type.  
`{ <KEY> = <TYPE>, <KEY> = <TYPE>, ... }`

For example: an object type of `object({ name=string, age=number })` would match a value like the following:

```
{  
  name = "John"  
  age   = 52  
}
```

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Types: Complex - Object

```
variable address_johndoe {  
  type=object({name=string,age=number})  
  default={name="john doe", age=33}  
}  
  
variable addressbook {  
  type=list(object({name=string,age=number}))  
  default=[ {name="wilson smith", age=34} ]  
}  
  
output addressbook { value = var.addressbook }  
output addressbook2 { value = concat(var.addressbook, [var.address_johndoe])}
```

Outputs:

```
addressbook = [ {  
  "age" = 34  
  "name" = "wilson smith"  
}, ]  
addressbook2 = [ {  
  "age" = 34  
  "name" = "wilson smith"  
}, {  
  "age" = 33  
  "name" = "john doe"  
}, ]
```

<https://www.terraform.io/docs/configuration/types.html>  
<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Types: Complex - Tuple

## Complex Types: Structural Types

- tuple: sequence of elements identified by consecutive whole numbers starting with zero, where each element has its own type.

a tuple type of tuple ([string, number, bool]) would match a value like the following:

```
["a", 15, true]
```

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Types: Complex - Tuple

```
variable atuple {  
  type=tuple([string,number,bool])  
  default=["john doe", 33, true]  
}  
  
output atuple { value = var.atuple }  
output atuple0 { value = var.atuple[0] }  
output atuple1 { value = var.atuple.1 }  
output atuple2 { value = var.atuple.2 }
```

Outputs:

```
atuple = [  
  "john doe",  
  33,  
  true,  
]  
atuple0 = john doe  
atuple1 = 33  
atuple2 = true
```

<https://www.terraform.io/docs/configuration/types.html>  
<https://www.terraform.io/docs/configuration/expressions.html>

# Terraform Control Structures



# Terraform Control Structures: loops

We can also “loop” over terraform lists:

```
count = length(some_list)
```

When the special attribute “count” is specified in a resource multiple resources are created.

We can access the current index in the resources with `count.index`

Note that elements begin with zero, not one!

# Terraform Control Structures: loops

We can use the ***count*** attribute to perform looping.  
For example if we have the following:

```
resource "aws_subnet" "vpc_subnets" {  
    # lines removed ...  
    count = length(var.vpc_subnet_cidr)  
  
    cidr_block = element(var.vpc_subnet_cidr, count.index)  
    tags = {  
        Name = "subnet-${count.index+1}"  
    }  
}
```

This will create several instances of this resource.  
Note the use of **count.index** to identify the loop number (starts at 0)

# Terraform if statements.

As of today (Nov 2019), Terraform doesn't natively support if/elif/else decision trees.

You have to be creative.

We can use count with a boolean variable and the “?” ternary operator

```
resource "aws_route53_health_check" "service_up" {  
    count = var.is_internal_alb ? 0 : 1  
    ...  
}
```

# Resource dependencies.

By studying the resource attributes used in interpolation expressions, Terraform automatically infers when one resource depends on another.

But there are cases where the dependency has to be specified – see example ...

# Resource dependencies

Notice the use of a resource id to specify the dependence

```
resource "aws_eip" "example" {  
    instance aws_instance.example.id  
}
```

```
resource "aws_instance" "example" {  
    ami = "ami-408c7f28"  
    instance_type = "t2.micro"  
    tags { Name = var.name }  
}
```

Terraform uses the dependency information to determine the correct order in which to create resources.

In this example, we create an *implicit dependency* so that Terraform knows to create the *aws\_instance* before the *aws\_eip*.

<https://github.com/hashicorp/terraform/blob/master/website/intro/getting-started/dependencies.html.md>

# Resource dependencies

Notice the use of a resource id to specify the dependence

```
resource "aws_s3_bucket" "example" { # S3 bucket for our application
  bucket = <unique bucket name>
  acl    = "private"
}
```

Sometimes there are dependencies between resources that are not visible to Terraform.

```
resource "aws_instance" "example" {
  ami           = "ami-2757f631"
  instance_type = "t2.micro"
```

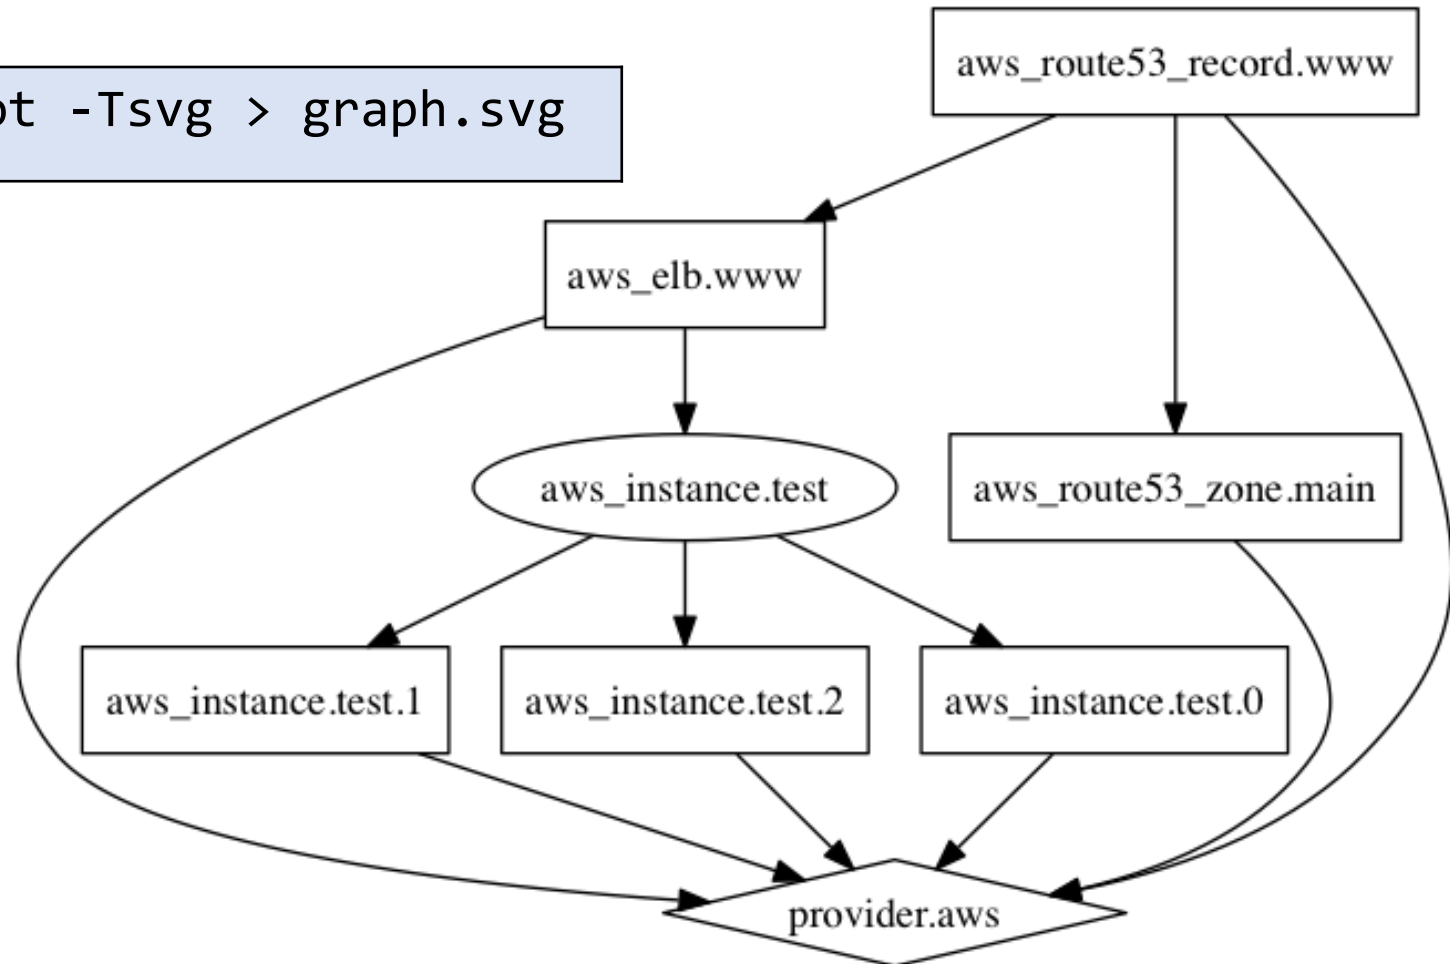
The *depends\_on* argument is accepted by any resource and accepts a list of resources to create explicit dependencies for

```
  # This EC2 instance to be created only after S3
  depends_on = [aws_s3_bucket.example] # explicit dependency
}
```

<https://github.com/hashicorp/terraform/blob/master/website/intro/getting-started/dependencies.html.md>

# Terraform graph

```
terraform graph | dot -Tsvg > graph.svg
```



<https://www.terraform.io/docs/commands/graph.html>

# Lab 3 – Terraform Lists

In this lab we will see how we can use variables of type list

- Setting of count to the length of the vpc\_subnet\_cidr list
- Setting individual cidr\_block attribute of aws\_instances to the appropriate cidr



# Quiz

What is special about the count variable ?

Can you loop 10 times with Terraform ?

What would happen if we change the value of count, then re-applied ?

# Terraform maps

Maps are key/value pairs.

Terraform supports declarations:

```
variable "mylist" {  
    type = "map"  
  
    default = {  
        "foo" = "bar",  
        "baz" = "blech"  
    }  
}
```

# Terraform functions

- Numeric: min, max
- String: chomp, join, split, replace, substr, upper, lower, ...
- List: length, sort, element, count (allows looping over list)
- Maps: lookup (looks up a map value based on a key)
- Filesystem: file, fileexists, dirname, ...
- Date/time: formatdate, timeadd, timestamp
- Hash/crypto: filemd5, filesha256, filesha512, md5, ...
- IP/network: cidrhost, cidrnetmask, cidrsubnet
- Type conversion: tobool, tolist, tomap, toset, tonumber, tostring

# Lab 4 – Terraform Maps

In this lab we will see how we can use variables of type map

- Using a map of regions to specify the ami image to use
  - Using a map of regions to availability zones

# Quiz

What is a possible use case of Maps ?

# Terraform Data Sources

# Terraform data sources

To get data directly from the cloud provider data sources.

```
data "aws_ami" "example" {  
  most_recent = true  
  
  owners = ["self"]  
  tags = {  
    Name      = "app-server"  
    Tested    = "true"  
  }  
}
```

<https://www.terraform.io/docs/providers/aws/>

Follow “Provider Data Sources” or “EC2”->“Data Sources”

# Terraform data sources

Note that we can declare “filters” to only get relevant data from the data source.

Note that “most\_recent” is a boolean.

Example:

“Find the latest available AMI that is tagged with Component = web”

```
data "aws_ami" "web" {  
  filter {  
    name      = "state"  
    values    = ["available"]  
  }  
  
  filter {  
    name      = "tag:Component"  
    values    = ["web"]  
  }  
  
  most_recent = true  
}
```



# Lab 5 – Terraform Data Sources

In this lab we will see how we can access Provider data sources

# Quiz

What is a possible use case of Data Sources ?

# Terraform State

# Terraform states

Terraform records the state locally in .tfstate json files by default:

```
> cat terraform.tfstate      # some lines removed
{ "version": 4,
  "terraform_version": "0.12.12", {
    "mode": "managed",
    "type": "aws_instance",
    "name": "example",
    "provider": "provider.aws",
    "instances": [ {
      "schema_version": 1,
      "attributes": {
        "ami": "ami-0e81aa4c57820bb57",
        "arn": "arn:aws:ec2:us-west-1:568285458700:instance/i-0f70aa9654b216df5",
        "associate_public_ip_address": true,
        "availability_zone": "us-west-1a",
        "cpu_core_count": 1,
```

# Terraform states

The previous state is saved as `terraform.tfstate.backup`

So after a “`terraform destroy`” you can still investigate the state of the previous configuration

The “`terraform show`” command allows to see a human (.tf) readable version of the state

# Coordinating Terraform states

When working in a team shared access to the Terraform state is necessary so you need

- Shared file system
- Locking
- Secrets

# Enabling remote state storage

The simplest “*remote backend*” solution when running on AWS is S3

It supports encryption, locking via DynamoDB, and versioning

It is possible to enable S3 storage to store terraform.tfstate files:

```
> terraform remote-config \  
  -backend=s3 \  
  -backend-config=bucket=my-s3-bucket  
  -backend-config=key=terraform.tfstate  
  -backend-config=encrypt=true  
  -backend-config=region=us-east-1
```

# Coordinating Terraform states

Hashicorp provides the *Atlas* / Terraform Cloud service which can store terraform.tfstate files.

It provides file locking, but it is expensive.

You can also create a Continuous Integration job manually with a tool like Jenkins.

Another good alternative is to use Terragrunt.



# Terraform remote state data source

Create a backend.tf containing:

```
backend = "s3" config = {  
    # Replace this with your bucket name!  
    bucket = "terraform-up-and-running-state"  
    key = "stage/data-stores/mysql/terraform.tfstate"  
    region = "us-east-2"  
}
```

Perform a “terraform init” to enable the remote storage backend

<https://blog.gruntwork.io/how-to-manage-terraform-state-28f5697e68fa> “How to manage Terraform State”  
<http://aws-cloud.guru/terraform-aws-backend-for-remote-state-files-with-s3-and-dynamodb/>

# Terraform remote state data source

The screenshot shows the AWS S3 console interface. At the top, the navigation bar includes the AWS logo, 'Services', 'Resource Groups', and user information 'mjbrightC'. The breadcrumb trail indicates the current location is 'Amazon S3 > 20191108-mjbright-s3'. Below the breadcrumb, there are four tabs: 'Overview', 'Properties', 'Permissions', and 'Management'. A search bar is present with the placeholder text 'Type a prefix and press Enter to search. Press ESC to clear.' Below the search bar, there are several action buttons: 'Upload', 'Create folder', 'Download', 'Actions', and 'Versions'. The 'Versions' section has 'Hide' and 'Show' buttons. The region is set to 'Asia Pacific (Singapore)'. The main content area shows a table with one file, 'terraform.tfstate', with a size of 3.3 KB and a storage class of 'Standard'. The file was last modified on Nov 8, 2019 at 6:49:42 AM GMT+0000.

aws Services Resource Groups mjbrightC Global Support

Amazon S3 > 20191108-mjbright-s3

Overview Properties Permissions Management

Q Type a prefix and press Enter to search. Press ESC to clear.

Upload Create folder Download Actions Versions Hide Show Asia Pacific (Singapore)

Viewing 1 to 1

<input type="checkbox"/> Name	Last modified	Size	Storage class
<input type="checkbox"/> terraform.tfstate	Nov 8, 2019 6:49:42 AM GMT+0000	3.3 KB	Standard

Viewing 1 to 1

# Terraform remote state data source

```
▼ {
  "version": 4,
  "terraform_version": "0.12.13",
  "serial": 0,
  "lineage": "052409a2-dfb2-0359-852f-971991c6c4c8",
  ▼ "outputs": {
    ▼ "dynamodb_table_name": {
      "value": "20191108-mjbright-dynamodb-lock",
      "type": "string"
    },
    ▼ "s3_bucket_arn": {
      "value": "arn:aws:s3:::20191108-mjbright-s3",
      "type": "string"
    }
  },
  ▼ "resources": [
    ▼ {
      "mode": "managed",
      "type": "aws_dynamodb_table",
      "name": "terraform_locks",
      "provider": "provider.aws",
      ▼ "instances": [
        ▼ {
          "schema_version": 1,
          ▼ "attributes": {
            "arn": "arn:aws:dynamodb:ap-southeast-1:568285458700:table/20191108-mjbright-dynamodb-lock",
            ▼ "attribute": [
              ▼ {
                "name": "LockID",
                "type": "S"
              }
            ]
          }
        }
      ]
    }
  ]
}
```

# Terraform State - Workspaces

Workspaces: allow to change context and create new resources isolated from resources in another context

“terraform workspace” allows to manage workspaces

- List
- Create
- Destroy

Warning: The workspace does not sufficiently protect the state when working in a multiple team member environment

<https://www.terraform.io/docs/state/workspaces.html>

# Lab 6 – Storing Persistent States

In this lab we will see how to store Terraform state information in a remote location such as an S3 bucket or a database

# Quiz

Why is it a good idea to store state in a remote backend ?

What are some examples of state backends ?

# Terraform Modules

# Terraform Modules

Terraform modules are directories that contain one or more terraform templates.

We can re-use these modules and templates, and subject them to version control.



# Terraform Modules

By convention we define three specific terraform templates in a module.

- vars.tf
- main.tf
- outputs.tf

# Vars.tf example

Specify module inputs in vars.tf:

```
variable "name" {  
    description = "The name of the EC2 instance"  
}  
  
variable "ami" {  
    description = "The AMI to run on the EC2 instance"  
}  
  
variable "port" {  
    description = "The port to listen on for HTTP requests"  
}
```

# main.tf example

Specify resources in main.tf:

```
resource "aws_instance" "example" {  
    ami           = var.ami  
    instance_type = "t2.micro"  
    user_data     = template_file.user_data.rendered  
    tags { Name = var.name }  
}
```

# Outputs.tf

Specify outputs in outputs.tf:

```
output "url" {  
    value = "http://${aws_instance.example.ip}:${var.port}"  
}
```

# Terraform modules

Note that terraform modules don't have "scope".  
Terraform modules don't share variables implicitly.

You must define inputs

- in the module vars.tf, and reference these from the root module

You must define outputs

- in the module resources.tf, so that other modules can use them.

Modules are like "functions" in other programming languages.

# Terraform modules

Terraform configs always have a 'root' module – the working directory where “*terraform apply*” is run

You 'get' modules with “*terraform get*” or “*terraform init*”.

Module locations can be specified with the source keyword.

Note that region is defined here as an input variable to module “vpc\_module”

```
module "vpc_module" {  
    source = "./modules/vpc"  
    region = var.region  
}
```

# Terraform modules

- Module sources can be:
  - Local files
  - Git repositories
  - URL's
  - Terraform registry locations.
  - Bitbucket
  - Mercurial repositories
  - Other ...

# Terraform modules

For variables to be accessible in the calling module, they must be defined as outputs in the called module

In the module: ./modules/vpc/output.tf


```
output "subnet_ids" {  
    value = aws_subnet.vpc_subnets.*.id  
}
```

In the “calling” module: modules/instances/main.tf

```
"aws_instance" "webserver" {  
    count      = length(module.vpc_module.aaz[var.region])  
    ami        = lookup(var.ami_instance, var.region)  
    subnet_id  = element(module.vpc_module.subnet_ids, count.index)
```




# Terraform Registry - Modules



## Find Terraform Modules

Use and learn from verified and community modules

### POPULAR MODULES



<https://registry.terraform.io/>

@mjbright CONSULTING

# Terraform Registry - Modules

Modules can be sourced from the Terraform public registry

Or from the Terraform Cloud private registry (paid service)

Or from any service that implements the registry API

<https://www.terraform.io/docs/registry/api.html>

<https://registry.terraform.io/>

# Lab 7 – Terraform Modules

In this lab we will see how we can split Terraform code into modules.  
More importantly we will see the advantages of modules for code management,  
testing, sharing

# Quiz

By convention what are the three files which make up a module ?

Modules are like functions, so how are their input/output specified ?

# Terraform Provisioners

# Terraform Provisioners

**Provisioners are not the same thing as providers !!**

Provisioners are typically used with Terraform to run remote commands on the created resource - provided that the resource supports it.

- Mainly used to run resource configuration management.
- Provisioners are added to resource definitions.

# Terraform Provisioners

They are provided for pragmatism and should generally be avoided as Terraform cannot predict their effects !

Can be used to model specific actions on the local machine or remote machine in order to prepare servers or other infrastructure objects for service.

A good use case would be to prepare (local hosts file) and execute **configuration management tools** which affect the internal state of resources

# Terraform Provisioners

Several in-built Provisioners exist:

- file
- habitat
- local-exec
- remote-exec
- chef
- puppet
- salt-masterless



# Terraform Provisioners

Local-exec performs an action on the local machine:

```
resource "aws_instance" "web" {  
  # ...  
  
  provisioner "local-exec" {  
    command = "echo ${self.private_ip} > file.txt"  
  }  
}
```

One use case might be to build up an `ansible_hosts` file

<https://www.terraform.io/docs/provisioners/local-exec.html>

<https://ilhicas.com/2019/08/17/Terraform-local-exec-run-always.html>

# Terraform Provisioners

Many provisioners are remote, and will require specific information to be supplied in order to connect.

<https://sdorsett.github.io/post/2018-12-26-using-local-exec-and-remote-exec-provisioners-with-terraform/>

# Terraform Provisioners

Example provisioner, “file”, used to copy files between local and remote systems:

```
provisioner "file" {  
  source      = "conf/myapp.conf"  
  destination = "/etc/myapp.conf"  
  
  connection {  
    type      = "ssh"  
    user      = "root"  
    password  = "${var.root_password}"  
  }  
}
```

# Terraform Provisioners

Example provisioners  
to copy then  
execute scripts  
on remote system:

```
provisioner "file" {  
    source      = "setup.sh"  
    destination = "setup.sh"  
}  
  
provisioner "remote-exec" {  
    inline = [ "chmod +x ./setup.sh", "./setup.sh" ]  
}  
  
connection {  
    type = "ssh"  
    host = self.public_ip  
    user = "ubuntu"  
    private_key = file(pathexpand("~/.ssh/id_rsa"))  
}
```

# Terraform Provisioners

Example provisioner to execute script on local system:

```
provisioner "local-exec" {  
  command = <<-EOF  
    echo "PRIVATE_IP=${aws_instance.example.private_ip}" >> ips.txt  
    echo "PUBLIC_IP=${aws_instance.example.public_ip}" >> ips.txt  
    echo "PUBLIC_DNS=${aws_instance.example.public_dns}" >> ips.txt  
  EOF  
}
```

This could be the basis for building up an `ansible_hosts` file

# Lab 8 – Setting up Web Servers

In this lab we will see how we can launch several Web Servers using Terraform

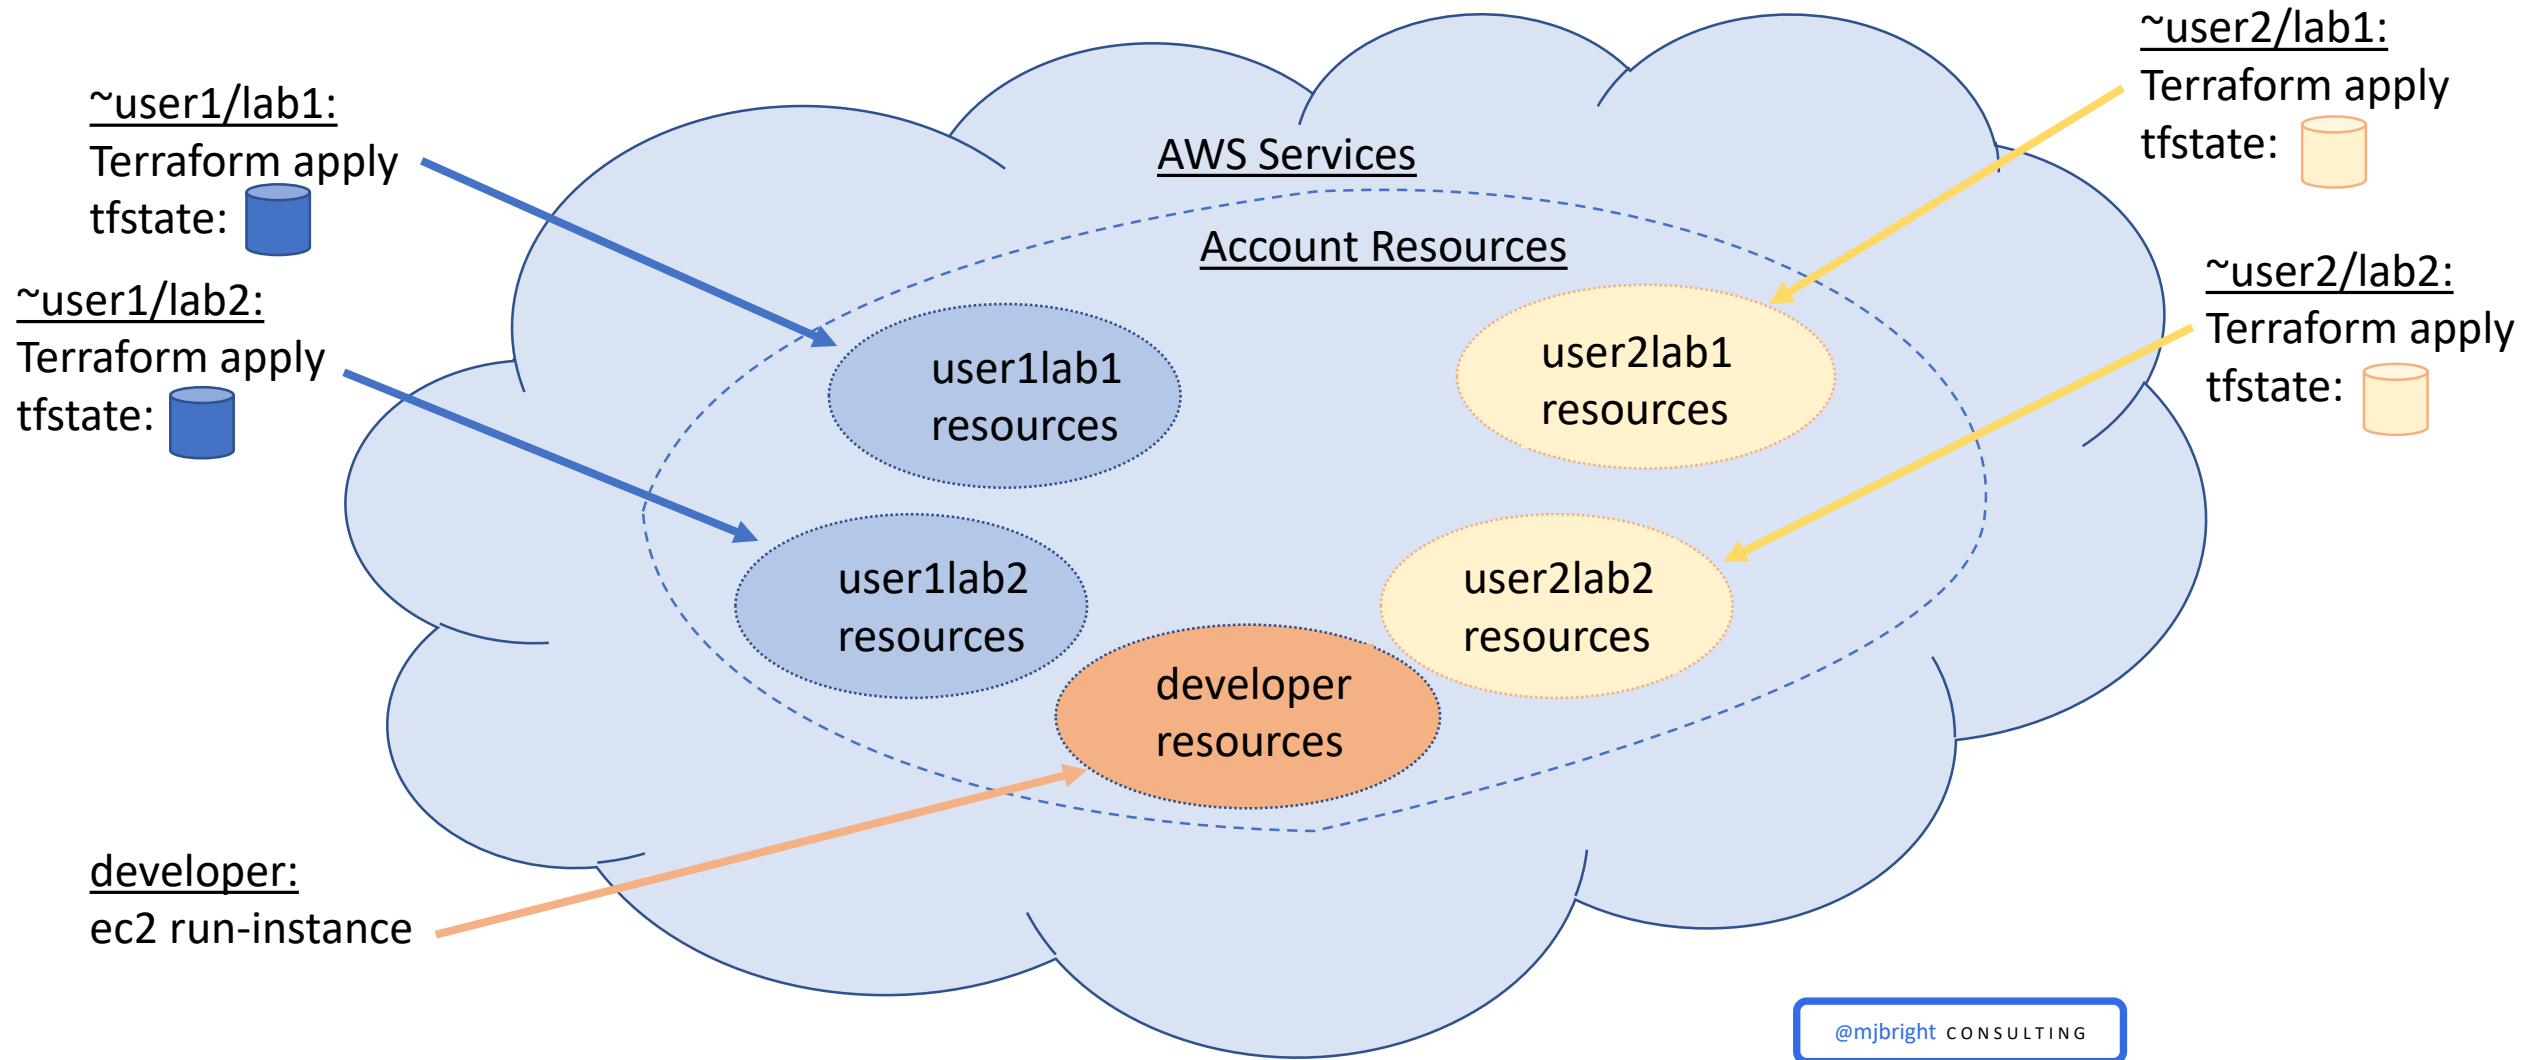
# Quiz

What are provisioners ?

Why should you avoid using them ?

When might you use them ?

# Terraform – State





# Terraform – Importing foreign resources

Terraform import

Terraforming

<https://github.com/dtan4/terraforming>

# Terraform – Import

**Terraform import:** Allows to import existing resources, from a provider

```
> terraform import aws_instance.example i-abcd1234
```

However, you first need to create a configuration file, e.g. `unmanaged.tf` with at least `ami` and `instance_type` specified:

```
resource "aws_instance" "example" {  
    ami = "ami-0e81aa4c57820bb57"  
    instance_type = "t2.micro"  
}
```

Note: It ignores any resources already created by Terraform

# Terraform – Import

Demo of Terraform import

# Terraforming – Import

**Terraforming:** Export AWS resources to Terraform style (tf, tfstate)

<http://terraforming.dtan4.net/>

<https://github.com/dtan4/terraforming>

e.g. export tf

```
> terraforming s3
resource "aws_s3_bucket" "hoge" {
    bucket = "hoge" acl = "private"
}
resource "aws_s3_bucket" "fuga" {
    bucket = "fuga" acl = "private"
}
```

e.g export tfstate

```
> terraforming s3 --tfstate
...
```

# Terraforming – Import

## **Terraforming:**

is written in ruby, so is a little more complicated to install

But it seems worth the effort if you need to dynamically manage “unmanaged” resources with Terraform

# Terraform – AWS Autoscaling & Loadbalancing

# Terraform – AWS ELB & ALB

AWS provide 3 types of Load balancer service

Elastic Load Balancer

- ELB: “Classic” Load-balancer, Layer4 routing, TCP, SSL/TLS, HTTP(S)

Application Load Balancer

- ALB: 2<sup>nd</sup>-gen Load-balancer, Layer7 routing, HTTP, HTTPS

Network Load Balancer

- NLB: 2<sup>nd</sup>-gen Load-balancer, Layer4 routing, TCP, UDP, TLS

<https://medium.com/@sahityamaruvada/setting-up-aws-network-load-balancer-with-terraform-0-12-b87e75992949>

For a comparison of AWS Load-balancers:

<https://aws.amazon.com/elasticloadbalancing/features/>

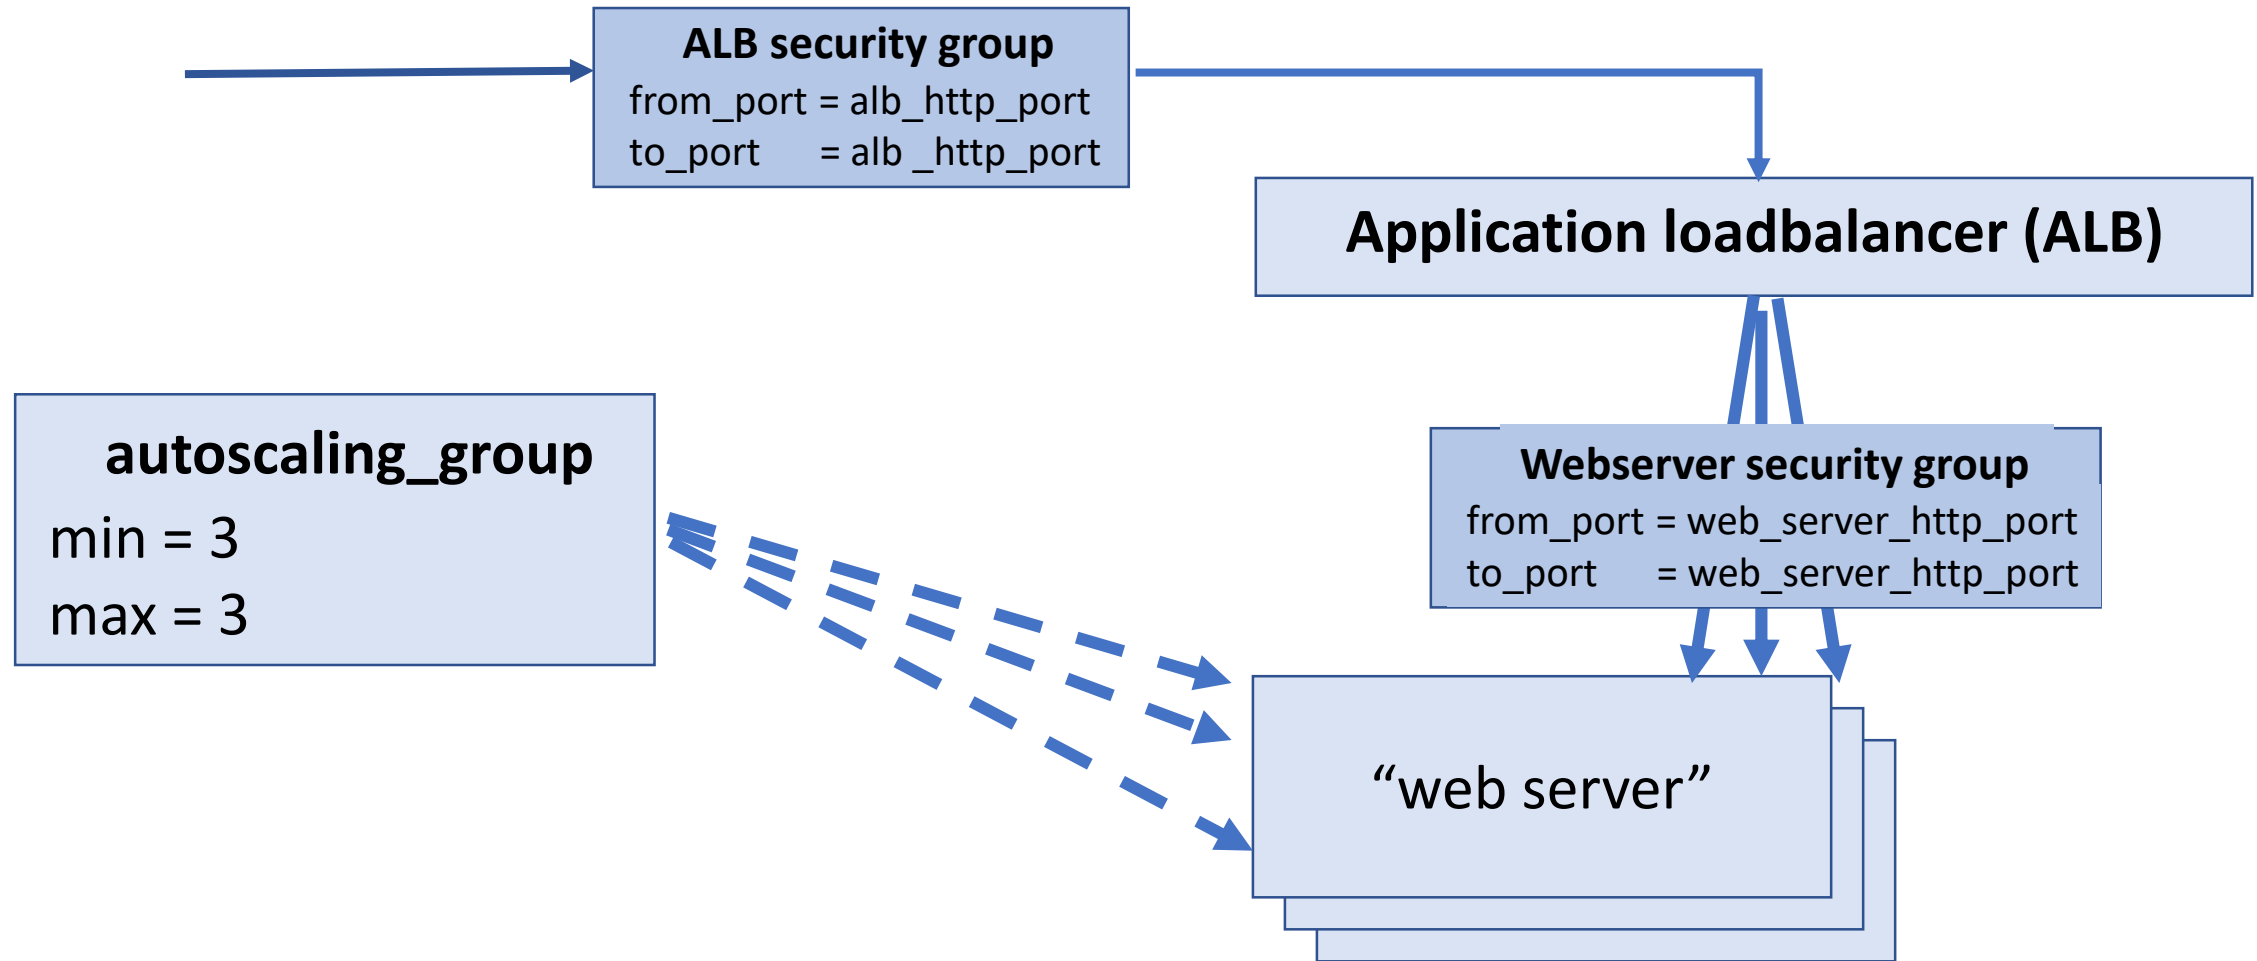
# Terraform – AWS Autoscaling & Loadbalancing

To demonstrate the use of AWS Autoscaling

- We will create an AutoScaling group (ASG) to scale the number of instances.
- We will associate this with an Application Load Balancer (ALB), which forwards traffic to the instances



# Terraform – AWS Autoscaling & Loadbalancing



# Terraform – AWS Autoscaling & Loadbalancing

## **autoscaling\_group**

min = 3

max = 3

The AutoScaling group (ASG) scales the number of instances.

The ASG uses a LaunchConfiguration to create the instances.  
In this simple example we fix at a constant 3 replicas

## **launch\_configuration**

image\_id = data.aws\_ami.ubuntu.id

instance\_type = "t2.micro"

```
lifecycle {  
  create_before_destroy = true  
}
```



# Terraform – AWS Autoscaling & Loadbalancing

```
resource "aws_autoscaling_group" "web_servers" {  
  name = aws_launch_configuration.web_servers.name  
  launch_configuration = aws_launch_configuration.web_servers.name  
  
  min_size      = 3  
  max_size      = 3  
  desired_capacity = 3  
  min_elb_capacity = 3  
  
  vpc_zone_identifier = data.aws_subnet_ids.default.ids # Deploy all the subnets (& AZs)  
  health_check_type = "ELB" # Register ASG instances in the ALB & use it's health check  
  target_group_arns = [aws_alb_target_group.web_servers.arn]  
  depends_on = ["aws_alb_listener.http"]  
}
```

# To support rolling deployments, must set in all dependencies also

```
lifecycle {  
  create_before_destroy = true  
}
```

# Terraform – AWS Autoscaling & Loadbalancing

```
resource "aws_launch_configuration" "web_servers" {  
  image_id      = data.aws_ami.ubuntu.id  
  instance_type = "t2.micro"  
  key_name      = var.key_name  
  
  lifecycle {  
    create_before_destroy = true  
  }  
  
  security_groups = [aws_security_group.web_server.id]  
  
  # To keep this example simple, we run a web server as a User Data script.  
  # In real-world usage, you would install in the AMI.  
  user_data = <<-EOF  
    #!/bin/bash  
    echo "${var.server_text} from $(hostname)" > index.html  
    nohup busybox httpd -f -p "${var.web_server_http_port}" &  
    EOF  
}
```

# Terraform – AWS Autoscaling & Loadbalancing

## ALB

### ALB Listener

```
port      = var.alb_http_port
protocol = "HTTP"
```

### ALB Listener Rule

```
type = "forward"
target_group_arn
```

### ALB Target Group

```
port      = var.web_http_port
Health_check
```

The ALB has a listener  
with a rule to listen on the web\_server\_http\_port.

The ALB forwards traffic to the ALB Target Group

```
lifecycle {
  create_before_destroy = true
}
```

“web server”

# Terraform – AWS Autoscaling & Loadbalancing

```
resource "aws_alb" "web_servers" {  
  name          = var.name  
  security_groups = [aws_security_group.alb.id]  
  subnets      = data.aws_subnet_ids.default.ids  
}
```

```
lifecycle {  
  create_before_destroy = true  
}
```

```
resource "aws_alb_listener" "http" {  
  load_balancer_arn = aws_alb.web_servers.arn  
  port              = var.alb_http_port  
  protocol          = "HTTP"  
  
  default_action {  
    type          = "forward"  
    target_group_arn = aws_alb_target_group.web_servers.arn  
  }  
}
```

# Terraform – AWS Autoscaling & Loadbalancing

```
resource "aws_alb_target_group" "web_servers" {  
  name      = var.name  
  port      = var.web_server_http_port  
  protocol  = "HTTP"  
  vpc_id    = data.aws_vpc.default.id
```

```
  lifecycle {  
    create_before_destroy = true  
  }
```

```
# Give existing connections 10 secs to complete before deregistering an instance ( default is 300 seconds )  
# In theory, ALB should deregister the instance when no open connections; in practice, it waits full delay.  
# If your requests are processed quickly, set to something lower (e.g. 10 seconds) to keep redeloys fast.  
deregistration_delay = 10
```

```
health_check {  
  path          = "/"  
  interval      = 15  
  healthy_threshold = 2  
  unhealthy_threshold = 2  
  timeout       = 5
```

# Terraform – AWS Autoscaling & Loadbalancing

```
resource "aws_alb_listener_rule" "send_all_to_web_servers" {  
  listener_arn = aws_alb_listener.http.arn  
  priority     = 100  
  
  action {  
    type          = "forward"  
    target_group_arn = aws_alb_target_group.web_servers.arn  
  }  
  
  condition {  
    field = "path-pattern"  
    values = ["*"]  
  }  
}
```



# Lab 9 – Scaling & loadbalancing web servers

In this lab we will see how to create the necessary resources to autoscale a group of Web Servers and to load-balance traffic to them

# Quiz

What resource is responsible to create the web server instances ?

Why do we specify “*create\_before\_destroy*” ?

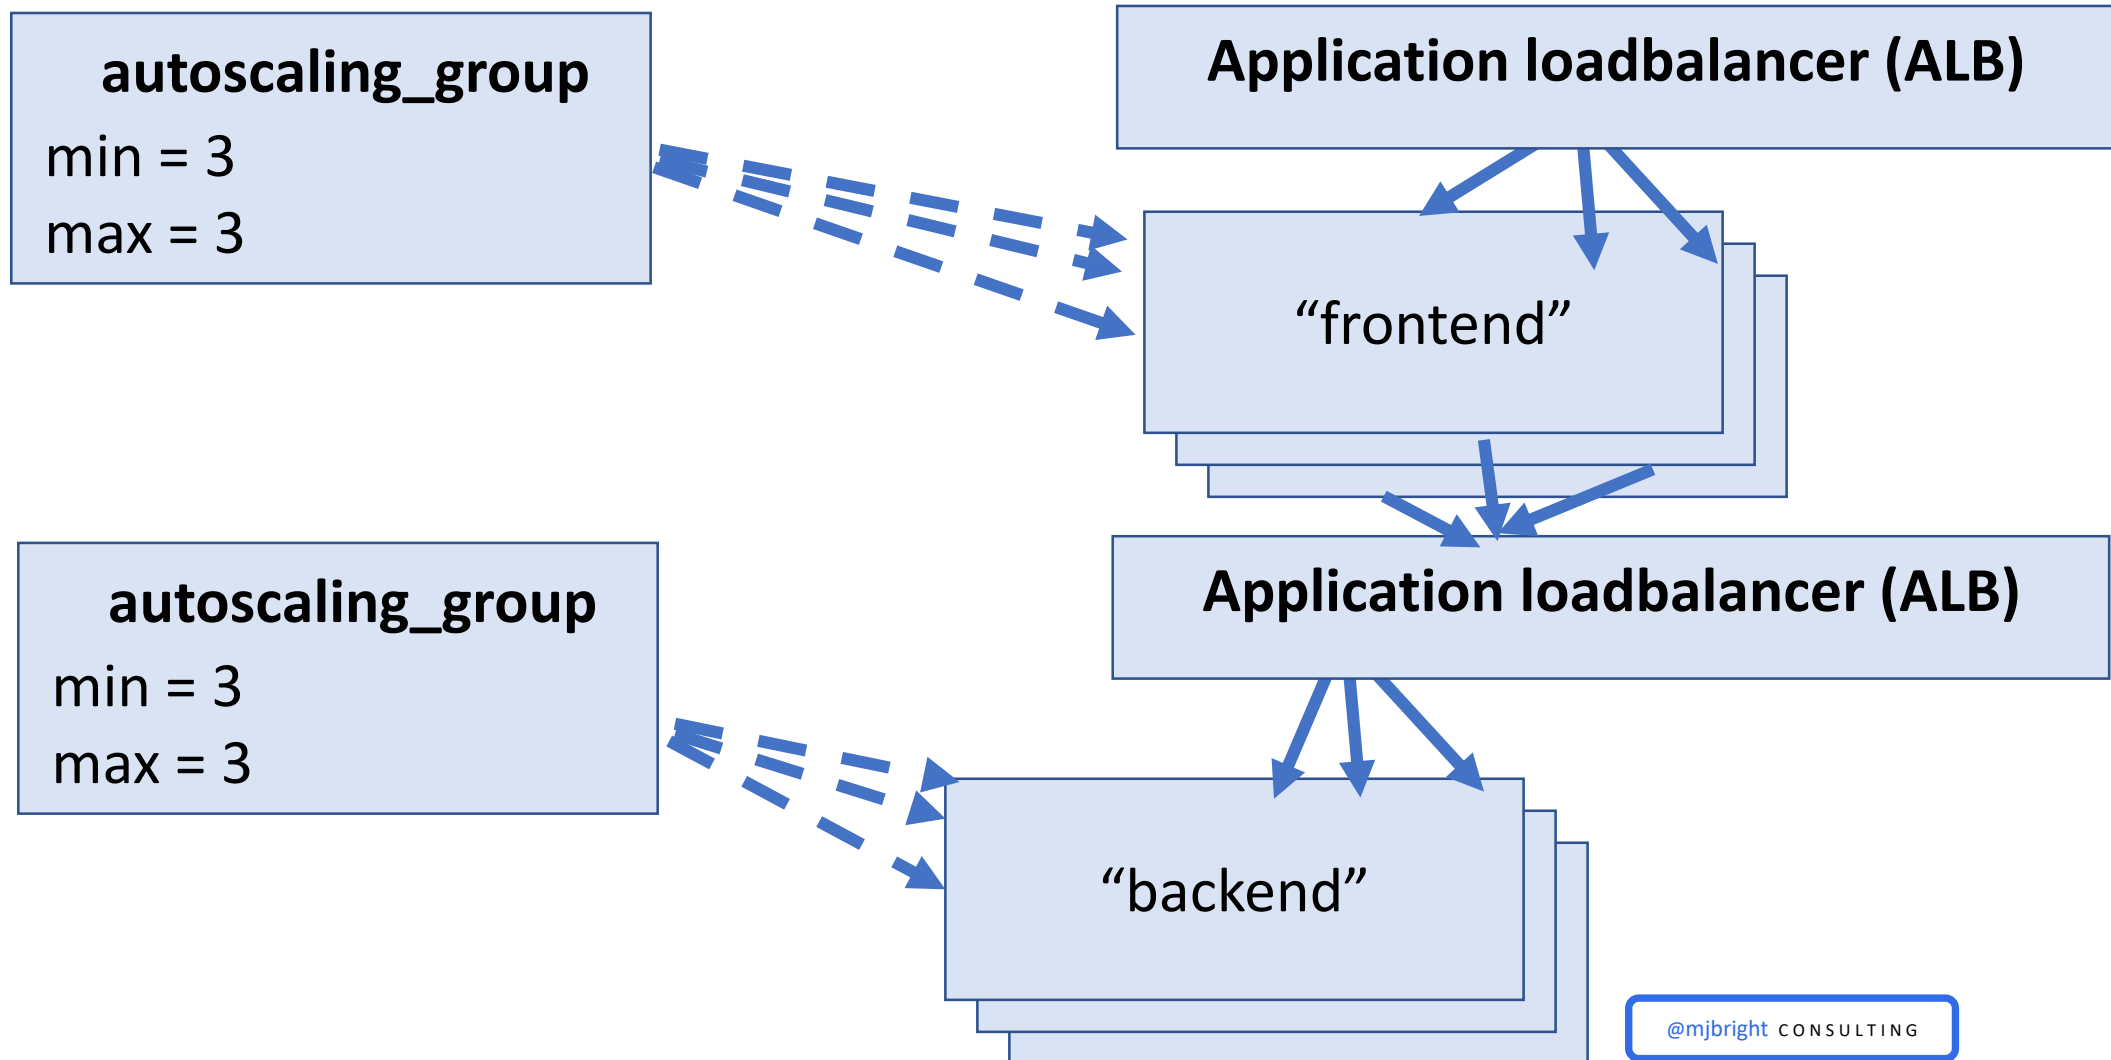
Why do we specify “*depends\_on = ["aws\_alb\_listener.http"]*” ?

How does the ALB TargetGroup verify that web\_servers are ready ?

# Terraform – AWS Autoscaling & Loadbalancing

We'll implement the scenario as a micro-service module facilitating the implementation of both backend & frontend load-balancing

# Terraform – AWS Autoscaling & Loadbalancing



# Lab 10 – Using a micro-service architecture

In this lab we will revisit the autoscaling & load-balancing across a group of Web Servers

This time we will refactor our code so that Autoscaling/Loadbalancing is implemented in a “*microservices*” module

This will allow us to autoscale/loadbalance backend and frontend resources independently

# Quiz

What are the 3 types of load balancer provided by AWS ?

What differentiates them ? On what basis would you choose ?

# Terraform – AWS VPCs

So far we have only used the default VPC – “Virtual Private Cloud”

Each region has a default VPC, but we can create up to 5 custom VPCs per region

Introduction to VPCs  
VPC & NAT  
EC2 instances in VPC

# Terraform – AWS VPCs

A VPC is a logically isolated virtual network where you can launch AWS resources

The virtual network has your choice of the following **AWS VPC components**:

- VPC CIDR Block
- subnet
- Internet gateway
- Route Table for Egress traffic
- Security group for Ingress/Egress traffic (stateful)
- Access Control Lists – firewall

Both IPv4 & IPv6 are available

<https://aws.amazon.com/vpc/>



# Terraform – AWS VPCs

The default VPC may already fit your needs

Reasons to set up a new VPC rather than using the default VPC

- some not-publicly reachable nodes - “default VPC” instances are reachable
- specific size of CIDR block
  - “default VPC” CIDR 172.31.0.0/16
- different subnet sizes
  - “default VPC” subnet per az of region CIDR 172.31.0.0/20

<https://aws.amazon.com/vpc/>

# Terraform – AWS VPCs

```
resource "aws_vpc" "vpc" {  
    cidr_block = var.cidr_block  
    instance_tenancy = "dedicated"  
    enable_dns_support = true  
    tags = {  
        Name = "${var.env}_vpc"  
        Env = var.env  
    }  
}  
  
resource "aws_subnet" "subnet" {  
    vpc_id = aws_vpc.vpc.id  
    cidr_block = var.subnet  
    map_public_ip_on_launch = "true"  
}
```

# Terraform – AWS VPCs

```
resource "aws_internet_gateway" "gw" {  
    vpc_id = aws_vpc.vpc.id  
}  
  
resource "aws_default_route_table" "route_table" {  
    default_route_table_id =  
aws_vpc.vpc.default_route_table_id  
  
    route {  
        cidr_block = "0.0.0.0/0"  
        gateway_id = aws_internet_gateway.gw.id  
    }  
}
```

# Terraform – AWS VPCs

```
resource "aws_instance" "VM-in-new-vpc" {  
    ami            = "ami-221ea342"  
    instance_type = "m3.medium"  
  
    subnet_id = var.subnet_id  
    vpc_security_group_ids = var.vpc_security_group_ids  
}
```

# Terraform – AWS EBS

Elastic Block Store (EBS) provides raw block-level storage attachable to Amazon EC2 instances

It is used by Amazon Relational Database Service (RDS)

Amazon EBS provides a range of options for storage performance and cost.

These options are divided into two major categories:

- SSD-backed storage for transactional workloads, such as DBs & boot volumes (IOPS)
- disk-backed storage for throughput intensive workloads, such as MapReduce & log processing (MB/s)

# Terraform – AWS EBS

```
resource "aws_ebs_volume" "ebs-volume-1" {
    availability_zone = us-west-1a"
    size = 20
    type = "gp2"
    tags {
        Name = "extra volume data"
    }
}

resource "aws_volume_attachment" "ebs-volume-1-attachment" {
    device_name = "/dev/xvdh"
    volume_id = "${aws_ebs_volume.ebs-volume-1.id}"
    instance_id = "${aws_instance.example.id}"
}
```

# Thank you !



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright

@mjbright CONSULTING

# Resources

Resource	URL
Website	<a href="https://www.terraform.io">https://www.terraform.io</a> <a href="https://www.terraform.io/docs">https://www.terraform.io/docs</a>
Downloads	<a href="https://www.terraform.io/downloads.html">https://www.terraform.io/downloads.html</a> <a href="https://github.com/hashicorp/terraform/releases">https://github.com/hashicorp/terraform/releases</a>
Learn Terraform	<a href="https://learn.hashicorp.com/terraform/getting-started/intro">https://learn.hashicorp.com/terraform/getting-started/intro</a> <a href="https://www.terraform.io/docs/glossary.html">https://www.terraform.io/docs/glossary.html</a>
Terraform Registry	<a href="https://registry.terraform.io/">https://registry.terraform.io/</a>



# Resources

Resource	URL
	<a href="https://www.terraform.io/docs/cloud/guides/recommended-practices/index.html">https://www.terraform.io/docs/cloud/guides/recommended-practices/index.html</a>
	<a href="https://www.terraform-best-practices.com">@antonbabenko</a>
	<a href="https://github.com/ozbillwang/terraform-best-practices">https://github.com/ozbillwang/terraform-best-practices</a>

# Resources

Resource	URL
	<a href="https://github.com/shuaibiyy/awesome-terraform"><u>https://github.com/shuaibiyy/awesome-terraform</u></a> <a href="https://github.com/kapil12345/awesome-terraform"><u>https://github.com/kapil12345/awesome-terraform</u></a>

# Resources

Resource	URL
Books	
	« Terraform Up & Running », O'Reilly,
	“The Terraform Book”, Turnbull Press
	“Terraform in Action”, Manning, Scott Winkler

# Resources

Resource	URL
	<a href="https://github.com/gruntwork-io/terraform-aws-couchbase.git">https://github.com/gruntwork-io/terraform-aws-couchbase.git</a>

# Thank you !



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright

@mjbright CONSULTING